

model_evaluation

December 17, 2025

```
[1]: import pandas as pd
import numpy as np
import warnings
from final_project.modeling import load_models, EVENT_WEIGHT
from final_project.data import read_data, split_data
from final_project.preprocessing import NUM_FEATURES, CAT_FEATURES, RESPONDER
from final_project.evaluation import evaluate_predictions, get_pred_summary, u
    ↪get_models_and_val_data
from final_project.plotting import plot_pred_vs_true, plot_day_predictions, u
    ↪plot_feature_relevance, plot_pdps

# Silence feature name warnings
warnings.filterwarnings(
    "ignore",
    message="X does not have valid feature names",
)
```

```
/opt/homebrew/Caskroom/miniconda/base/envs/final-project/lib/python3.10/site-
packages/dalex/_global_checks.py:1: UserWarning: pkg_resources is deprecated as
an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The
pkg_resources package is slated for removal as early as 2025-11-30. Refrain from
using this package or pin to Setuptools<81.
    import pkg_resources
```

0.1 Load Models

```
[2]: glm_raw, lgbm_raw, X_val_raw, y_val_raw = get_models_and_val_data("clean_data")
glm_clip, lgbm_clip, X_val_clip, y_val_clip = u
    ↪get_models_and_val_data("clean_data_clipped")
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000821 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 4702
[LightGBM] [Info] Number of data points in the train set: 21452, number of used
features: 74
[LightGBM] [Info] Start training from score -0.789211
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
```

```
testing was 0.000656 seconds.  
You can set `force_row_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force_col_wise=true`.  
[LightGBM] [Info] Total Bins 4702  
[LightGBM] [Info] Number of data points in the train set: 21452, number of used  
features: 74  
[LightGBM] [Info] Start training from score -0.789211
```

```
[3]: df_pred_raw = get_pred_summary(glm_raw, lgbm_raw, X_val_raw, y_val_raw)  
df_pred_clip = get_pred_summary(glm_clip, lgbm_clip, X_val_clip, y_val_clip)
```

0.2 Evaluate Models

Added baseline (past_50m_span_ewm_vol) to compare to models.

```
[4]: # Evaluate glm  
glm_eval_raw = evaluate_predictions(  
    df_pred_raw["y_true"], df_pred_raw["glm_y_pred"],  
    df_pred_raw["weight"]  
)  
print("Unclipped:\n")  
print(glm_eval_raw)  
glm_eval_clip = evaluate_predictions(  
    df_pred_clip["y_true"], df_pred_clip["glm_y_pred"],  
    df_pred_clip["weight"]  
)  
print("\nClipped:\n")  
print(glm_eval_clip)
```

Unclipped:

```
bias          0.023247  
gamma_deviance 0.124871  
mae           0.096982  
rmse          0.188458  
dtype: float64
```

Clipped:

```
bias          0.023018  
gamma_deviance 0.124766  
mae           0.096265  
rmse          0.168332  
dtype: float64
```

```
[5]: # Evaluate lgbm  
lgbm_eval_raw = evaluate_predictions(  
    df_pred_raw["y_true"], df_pred_raw["lgbm_y_pred"],
```

```

        df_pred_raw["weight"]
    )
print("Unclipped:\n")
print(glm_eval_raw)
lgbm_eval_clip = evaluate_predictions(
    df_pred_clip["y_true"], df_pred_clip["lgbm_y_pred"],
    df_pred_clip["weight"]
)
print("\nClipped:\n")
print(lgbm_eval_clip)

```

Uncropped:

```

bias           0.023247
gamma_deviance 0.124871
mae            0.096982
rmse           0.188458
dtype: float64

```

Clipped:

```

bias           0.007185
gamma_deviance 0.106868
mae            0.086975
rmse           0.156047
dtype: float64

```

```
[6]: # Evaluate baseline
base_eval_raw = evaluate_predictions(
    df_pred_raw["y_true"], df_pred_raw["baseline_y_pred"],
    df_pred_raw["weight"]
)
print("Unclipped:\n")
print(base_eval_raw)
base_eval_clip = evaluate_predictions(
    df_pred_clip["y_true"], df_pred_clip["baseline_y_pred"],
    df_pred_clip["weight"]
)
print("\nClipped:\n")
print(base_eval_clip)
```

Uncropped:

```

bias           -0.000721
gamma_deviance 0.138242
mae            0.097388
rmse           0.179108
dtype: float64

```

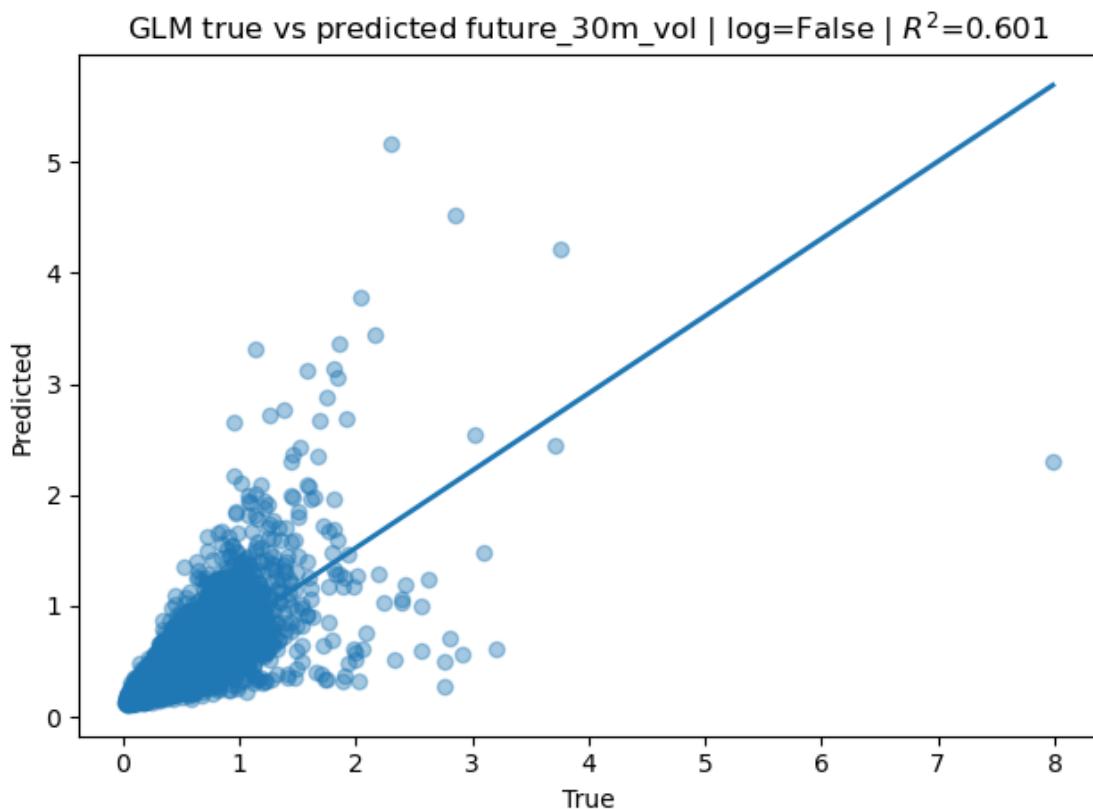
Clipped:

```
bias           -0.001047
gamma_deviance 0.138200
mae            0.097066
rmse           0.175826
dtype: float64
```

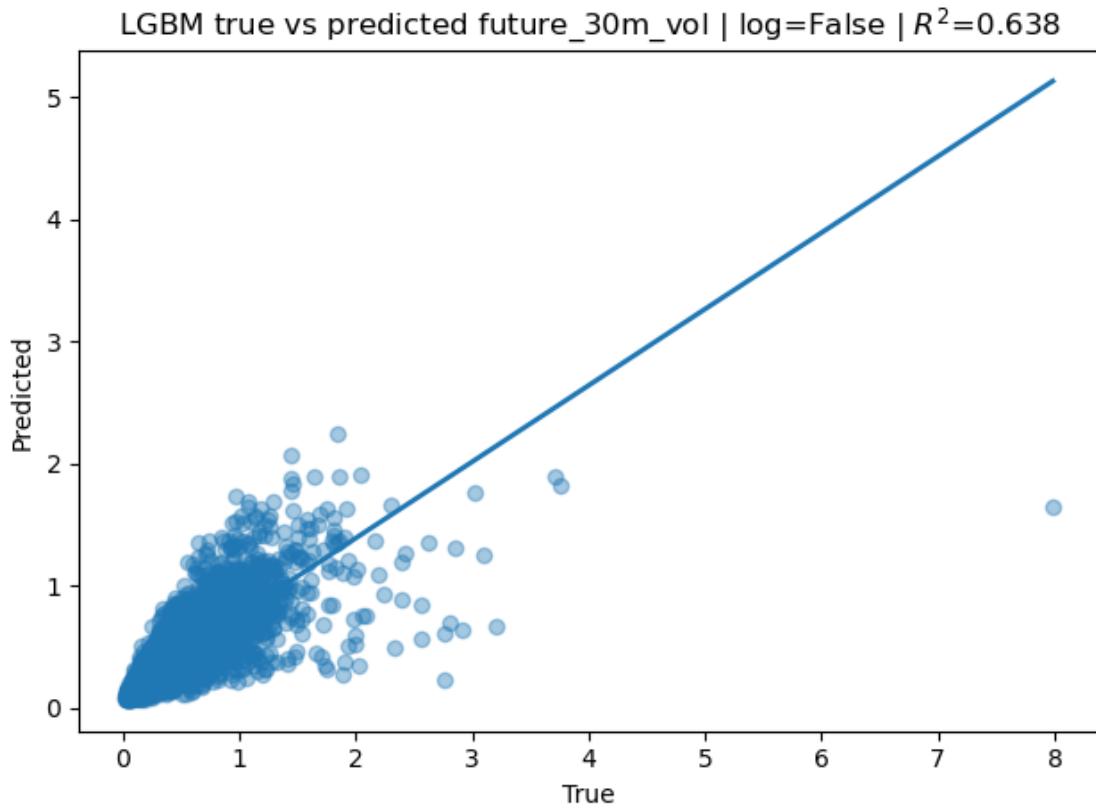
Looks like the GBT outperformed the GLM on all measures! And clipped data is significantly better than unclipped.

Predicted vs. Actual

```
[7]: # Plot pred vs. true for glm
fig = plot_pred_vs_true(df_pred_clip, "glm")
```

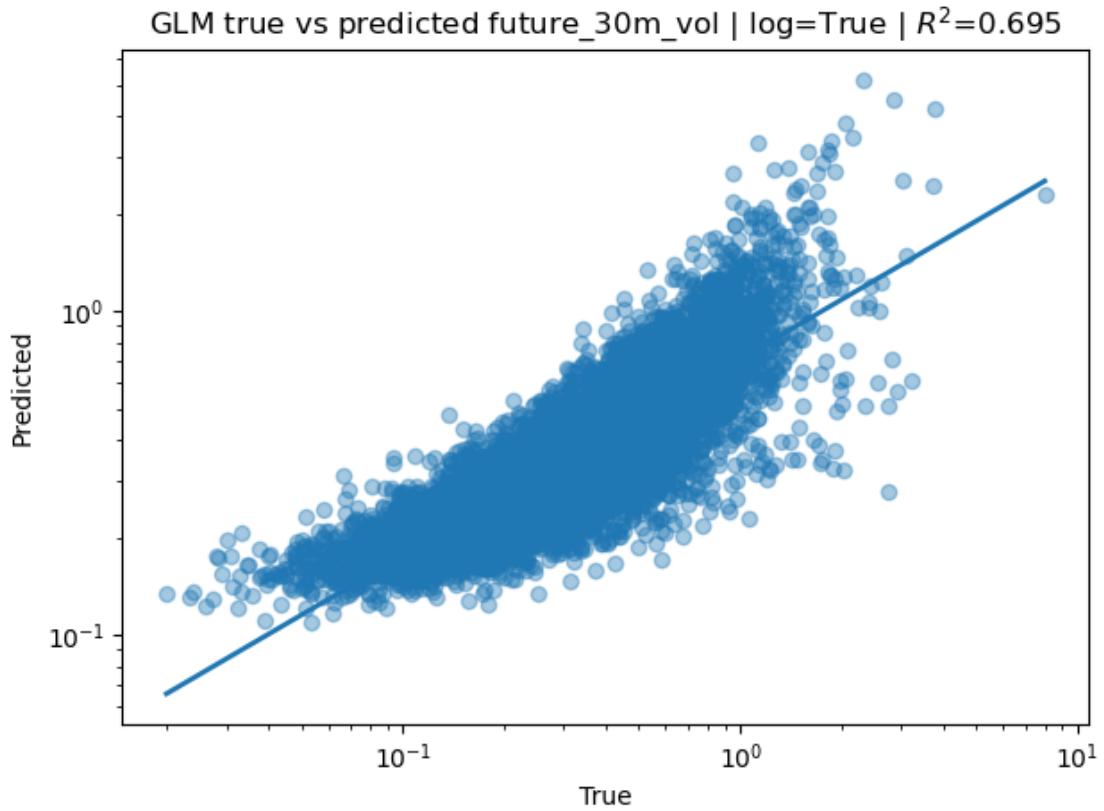


```
[8]: # Plot pred vs. true for lgbm
fig = plot_pred_vs_true(df_pred_clip, "lgbm")
```



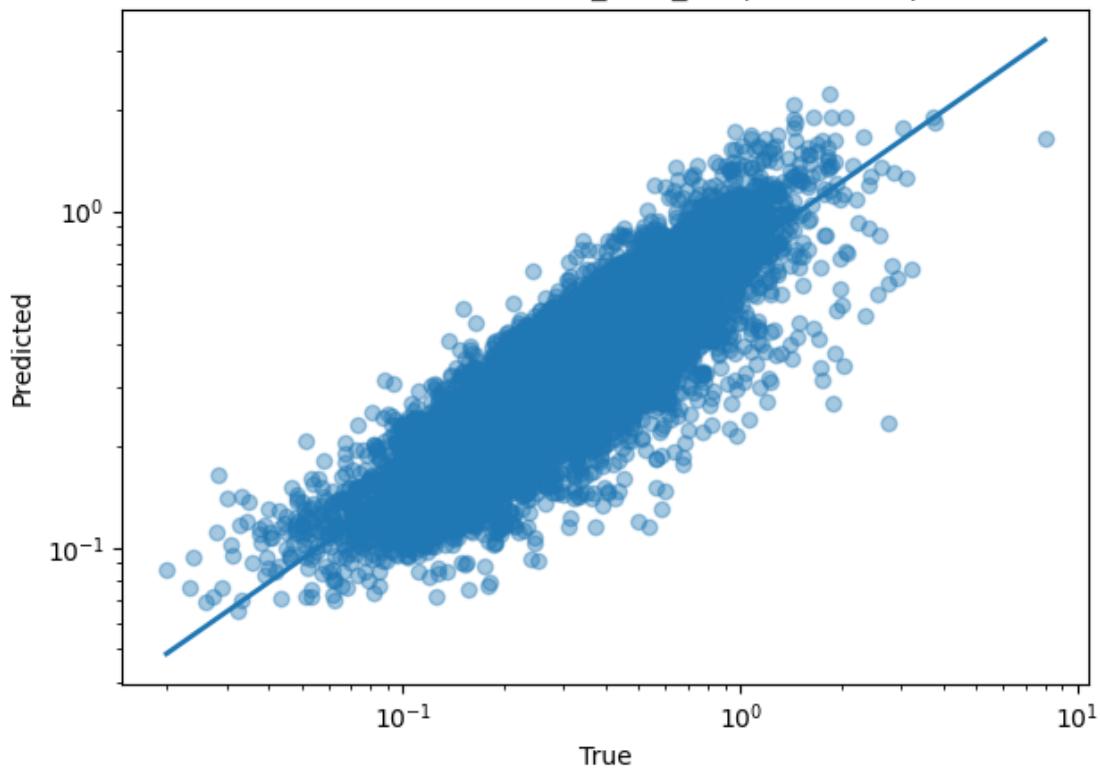
This should look better with logs!

```
[9]: # Plot pred vs. true for glm, log axes
fig = plot_pred_vs_true(df_pred_clip, "glm", log=True)
```



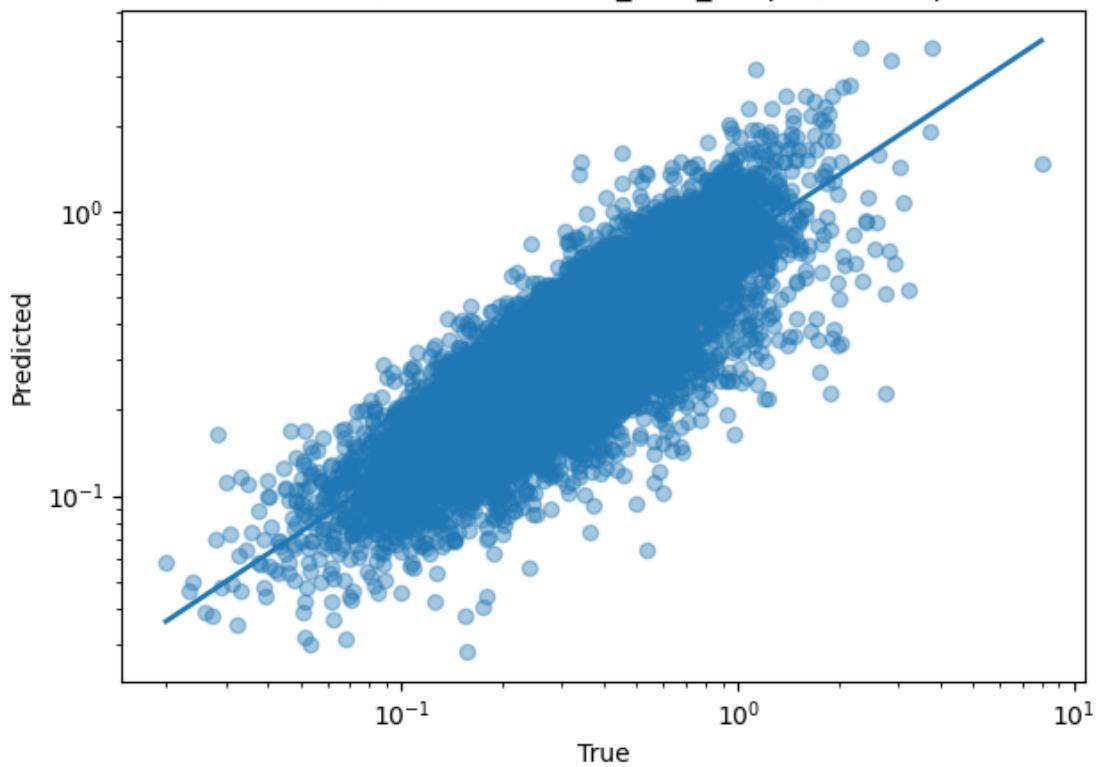
```
[10]: # Plot pred vs. true for lgbm, log axes
fig = plot_pred_vs_true(df_pred_clip, "lgbm", log=True)
```

LGBM true vs predicted future_30m_vol | log=True | $R^2=0.730$



```
[11]: # Plot pred vs. true for baseline, log axes
fig = plot_pred_vs_true(df_pred_clip, "baseline", log=True)
```

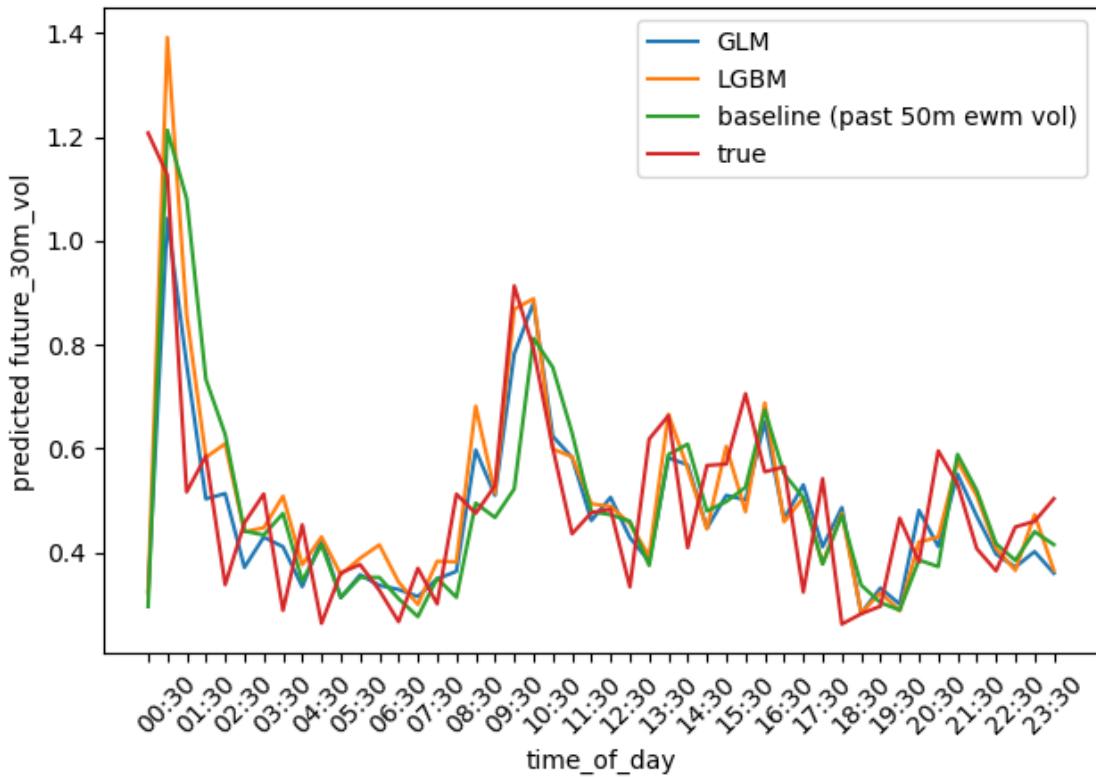
BASELINE true vs predicted future_30m_vol | log=True | $R^2=0.678$

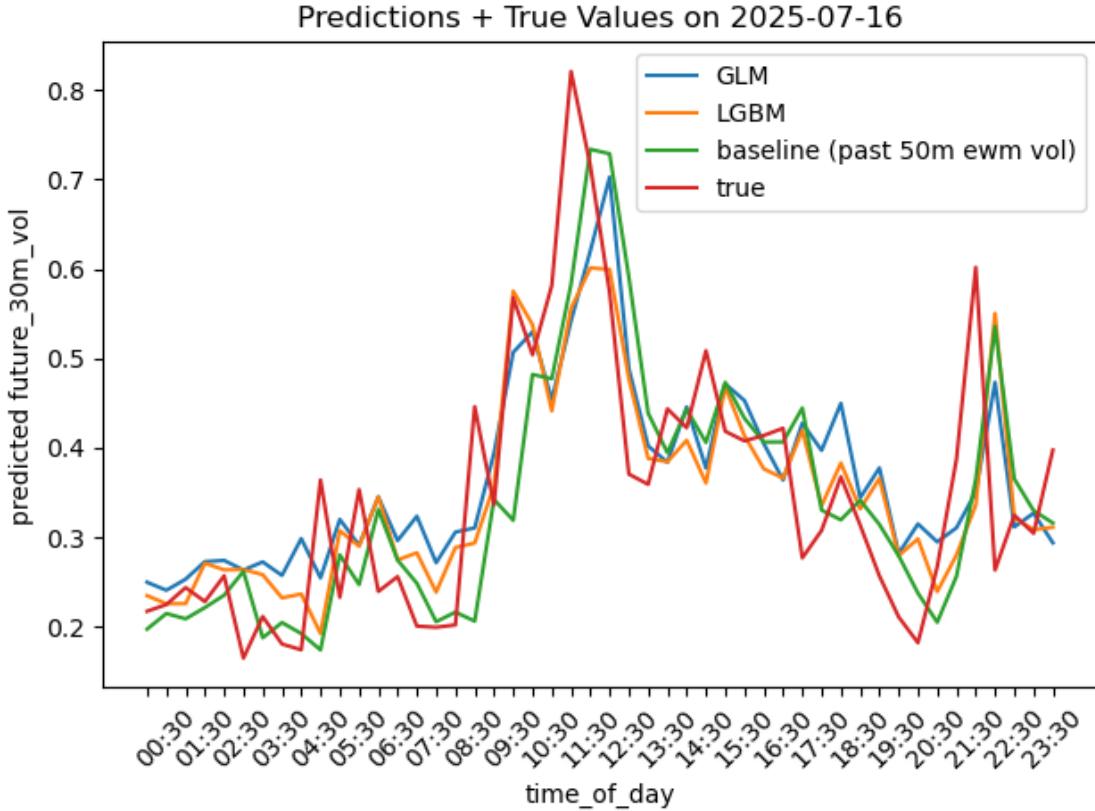


Look at predicted value throughout the day.

```
[12]: fig = plot_day_predictions(df_pred_clip, "2025-10-30")
fig = plot_day_predictions(df_pred_clip, "2025-07-16")
```

Predictions + True Values on 2025-10-30





0.3 Feature Relevance and PDPs

```
[13]: # Plot glm features
glm_top_5 = plot_feature_relevance(glm_clip, X_val_clip, y_val_clip)
```

Preparation of a new explainer is initiated

```
-> data : 14301 rows 29 cols
-> target variable : Parameter 'y' was a pandas.Series. Converted to a
numpy.ndarray.
-> target variable : 14301 values
-> model_class : glum._glm.GeneralizedLinearRegressor (default)
-> label : Not specified, model's class short name will be used.
(default)
-> predict function : <function yhat_default at 0x3021b2e60> will be used
(default)
-> predict function : Accepts only pandas.DataFrame, numpy.ndarray causes
problems.
-> predicted values : min = 0.109, mean = 0.38, max = 5.16
-> model type : regression will be used
-> residual function : difference between y and yhat (default)
```

```
-> residuals      : min = -2.86, mean = -0.023, max = 5.69
-> model_info     : package sklearn
```

A new explainer has been created!

```
[14]: # Plot lgmb features
lgbm_top_5 = plot_feature_relevance(lgbm_clip, X_val_clip, y_val_clip)
```

Preparation of a new explainer is initiated

```
-> data          : 14301 rows 29 cols
-> target variable : Parameter 'y' was a pandas.Series. Converted to a
numpy.ndarray.
-> target variable : 14301 values
-> model_class    : lightgbm.sklearn.LGBMRegressor (default)
-> label          : Not specified, model's class short name will be used.
(default)
-> predict function : <function yhat_default at 0x3021b2e60> will be used
(default)
-> predict function : Accepts only pandas.DataFrame, numpy.ndarray causes
problems.
-> predicted values : min = 0.0648, mean = 0.364, max = 2.24
-> model type      : regression will be used
-> residual function : difference between y and yhat (default)
-> residuals       : min = -0.774, mean = -0.00719, max = 6.34
-> model_info       : package sklearn
```

A new explainer has been created!

```
[15]: plot_pdps(glm_clip, X_val_clip, y_val_clip, n_top=5)
```

Preparation of a new explainer is initiated

```
-> data          : 14301 rows 29 cols
-> target variable : Parameter 'y' was a pandas.Series. Converted to a
numpy.ndarray.
-> target variable : 14301 values
-> model_class    : glum._glm.GeneralizedLinearRegressor (default)
-> label          : Not specified, model's class short name will be used.
(default)
-> predict function : <function yhat_default at 0x3021b2e60> will be used
(default)
-> predict function : Accepts only pandas.DataFrame, numpy.ndarray causes
problems.
-> predicted values : min = 0.109, mean = 0.38, max = 5.16
-> model type      : regression will be used
-> residual function : difference between y and yhat (default)
-> residuals       : min = -2.86, mean = -0.023, max = 5.69
-> model_info       : package sklearn
```

```

A new explainer has been created!

Calculating ceteris paribus: 100%| 4/4 [00:00<00:00, 37.37it/s]
Calculating ceteris paribus: 100%| 1/1 [00:00<00:00, 58.76it/s]

[16]: plot_pdps(lgbm_clip, X_val_clip, y_val_clip, n_top=5)

Preparation of a new explainer is initiated

-> data : 14301 rows 29 cols
-> target variable : Parameter 'y' was a pandas.Series. Converted to a
numpy.ndarray.
-> target variable : 14301 values
-> model_class : lightgbm.sklearn.LGBMRegressor (default)
-> label : Not specified, model's class short name will be used.
(default)
-> predict function : <function yhat_default at 0x3021b2e60> will be used
(default)
-> predict function : Accepts only pandas.DataFrame, numpy.ndarray causes
problems.
-> predicted values : min = 0.0648, mean = 0.364, max = 2.24
-> model type : regression will be used
-> residual function : difference between y and yhat (default)
-> residuals : min = -0.774, mean = -0.00719, max = 6.34
-> model_info : package sklearn

A new explainer has been created!

Calculating ceteris paribus: 75%| 3/4 [00:00<00:00,
9.01it/s]/opt/homebrew/Caskroom/miniconda/base/envs/final-
project/lib/python3.10/site-
packages/dalex/predict_explanations/_ceteris_paribus/utils.py:100:
FutureWarning:

Setting an item of incompatible dtype is deprecated and will raise in a future
error of pandas. Value '[0. 0.01 0.02 ... 0.98 0.99 1.]' has dtype
incompatible with int64, please explicitly cast to a compatible dtype first.

Calculating ceteris paribus: 100%| 4/4 [00:00<00:00, 8.98it/s]
Calculating ceteris paribus: 100%| 1/1 [00:00<00:00, 16.24it/s]

```